CAPSTONE PROJECT BY TEAM 3

A DATA SCIENCE APPROACH TO FORECAST SHORT-TERM ELECTRICITY DEMAND IN NSW

Karim Adham (z5468227), Rishantha Rajakaruna (z5441528), Rushmila Islam (z5456038)

School of Mathematics and Statistics
UNSW Sydney

September 2024

Plagiarism statement

I declare that this thesis is my own work, except where acknowledged, and has not been submitted for academic credit elsewhere.

I acknowledge that the assessor of this thesis may, for the purpose of assessing it:
- Reproduce it and provide a copy to another member of the University; and/or,
- Communicate a copy of it to a plagiarism checking service (which may then retain a copy of it on its database for the purpose of future plagiarism checking).

I certify that I have read and understood the University Rules in respect of Student Academic Misconduct, and am aware of any potential plagiarism penalties which may apply.

By signing this declaration I am agreeing to the statements and conditions above.

Signed: _____   Date: _____

Signed: _____   Date: _____

Signed: _____   Date: _____

Acknowledgements

TBW. . .

02/10/2024.

Abstract

TBW . . .

Contents

CHAPTER 1

Introduction

This R Markdown template can be used for the ZZSC9020 course report. You can incorporate R [R Core Team(2017)] chunks and Python chunks that will be run on the fly. You can incorporate LaTeXcommands.

TBW

Literature Review

2.1 Background

Electricity demand forecasting is crucial for ensuring an efficient, reliable and cost-effective power operation systems. Accurate forecasts enable grid operators to balance supply and demand in real-time, preventing costly overproduction or dangerous shortages that could lead to blackouts. It also helps to optimize the scheduling of power generation for different sectors, reducing operational costs and enhancing system reliability. As renewable energy sources like solar and wind become more integrated, forecasting plays a vital role in managing fluctuations in supply. Additionally, in deregulated markets, precise demand predictions allow energy providers to make informed trading decisions, ultimately benefiting both suppliers and consumers with more stable and affordable electricity.

NSW is largely self-sufficient in relation to electricity supply according to [www.soe.epa.nsw.g meeting most of its local demand through state generation. The remaining electricity is purchased from other states through the National Electricity Market (NEM). Electricity imports enable NSW to manage its supply at the lowest cost to consumers. In 2019–20, renewable energy sources accounted for approximately 19% of the state's total electricity generation, a significant increase from just over four times the amount in 2008–09.

According to [NSW EPA(2021)] electricity demand from the NSW grid is projected to experience a slight decline over the next five years before rising more. As stated in the ariticle the main reason for the decline in energy consumption is the lower consumption by the NSW industrial sector, particularly the manufacturing industry, since 2012–13. Industrial energy consumption saw a marked decline after the closure of the Kurnell and Clyde petroleum refineries in October 2014, with few significant changes observed since then. The decrease is expected to be driven by population growth being balanced by improvements in the energy efficiency of appliances and machinery. In addition, the increasing adoption of rooftop solar panels and battery storage systems is anticipated to further lower residential demand on the electricity grid. However, beyond the five-year mark, consumption is forecasted to increase as electric vehicle charging and broader electrification begin to significantly impact electricity demand. According to NSW Climate and Energy Action. (n.d.) the share of solar and wind in NSW's energy mix has more than DOUBLED from 5% in 2015 to 12% in 2019. In the Figure 2.1 we can visualize the rapid growth of electricity generation using renewable sources.

Figure 2.1: Renewable fuel sources [Source: Derived from Department of the Environment and Energy, Australian Energy Statistics, Table O, June 2021]

As stated in [www.soe.epa.nsw.gov.au()] the driving factors of electricity consumption and demand forecasts can be split into two different types: structural drivers and random drivers. Structural drivers such as population, economic growth, electricity price, technology adoption, energy efficiency can be estimated based on past trends and expert judgement. On the other hand, random drivers, such as weather, customer behavior etc, which can be modeled as probability distributions. There are many factors that drive consumers to make similar choices regarding electricity consumption at the same time for example during work and school schedules. And the demand is different during weekdays, public holidays, weekends, due to weather the use of heating and cooling appliances, and many other societal factors, such as whether the beach is pleasant, or the occurrence of retail promotions.

Load forecasting is usually concerned with the prediction of hourly, daily, weekly, and annual values of the system demand and peak demand. Such forecasts are sometimes categorized as short-term, medium-term and long-term forecasts, depending on the time horizon. In terms of forecasting outputs, load forecasts can also be categorized as point forecasts (i.e., forecasts of the mean or median of the future demand distribution), and density forecasts (providing estimates of the full probability distributions of the possible future values of the demand).

2.2 Factors Affecting Load Forecasting

In addition to the basic components of electricity demand forecasting, several external factors play a significant role in refining predictions and improving accuracy. Typically, electricity consumption is influenced by two factors—1) structural like

population growth, economic condition, electricity price, etc; and 2) non-structural or random like weather condition e.g., air temperature, extreme heat or cold, bushfire, flood, etc., building type e.g., multi-storey or free-standing house, and adoption of electric vehicles and contributions of renewable energy in the power grid, etc.

Temperature is a crucial factor, as it directly influences electricity consumption patterns. Extreme temperatures, whether very hot or cold, can lead to higher demand for heating or cooling, respectively. Forecasts often need to account for temperature variations to predict demand more accurately, especially during seasonal extremes or unusual weather patterns.


Figure 2.2: A caption1

As the graph in Figure 2.2 shows the temperature in Australia is rising and forecast show that it will continue to rise. According to [Adapt NSW(2024)] the climate of NSW is changing, with 6 of the 10 warmest years on record occurring in the past 10 years. The warmest year on record for both average temperature and maximum temperature in NSW was 2019, when average temperature was 1.2°C above the 1990–2009 average. Across NSW, average temperatures will continue to increase throughout this century. By 2090, average temperature is projected to rise by around 1.3°C under a low emissions scenario and around 4.0°C under a high emissions scenario. So temperature is key factor while determining the electricity demand. The usage of heating in colder nights and cooling system during the hot days directly impacts the demand of electricity.

Holidays and weekdays are also important key factors in load forecasting. On holidays, electricity consumption patterns can differ significantly from regular weekdays due to changes in work routines and social activities. For instance, public

holidays might see a decrease in commercial and industrial electricity use, while residential usage could increase due to family gatherings and home activities. Understanding these variations helps in adjusting forecasts to better reflect the actual demand patterns during such periods.

Overall, incorporating temperature data and considering holiday effects are essential for creating more accurate and reliable electricity demand forecasts, which in turn support better decision-making for power system management and operational planning.

2.3 Electricity Demand Forecast Methodologies

Electricity demand or load forecasting is a well-known problem to predict the future load based on historic information. Researchers used both statistical and probabilistic methods to achieve high accuracy with low errors to find the best model. They also considered this problem as both deterministic and stochastic to include the influence of the external factors that sways the models behaviour at different time horizons. As mentioned in [Ahmad et al.(2022)Ahmad, Ghadi, Adnan and Ali] load forecasting on the basis of time horizon, can be classified into four forecasting groups – VSLTF (very short-term (minutes to hours)), STLF (short-term (mitues to days)), MTLF (medium-term (days to months)), and LTF (long-term (months to years)).

There are very few methods for VSTLF including genetic algorithm, autoregressive moving average models and artificial neural network. STLF is utilized for time hardly from minutes to hours. STLF is the important source of information for daily operations and it is important for system operations [36]. Researchers are taking more interest to design predictive models because STLF can be used to approximate the long time load. It is essential to have accurate predict knowledge of affecting factors to improve short term model. The relation between demand of load and factors is the basic purpose to look for because instantaneous demand may be different. For duration of days to months, usually MTLF is used for load forecasting [37]. It becomes popular in peak summer or winter. For load duration from few weeks to many years, LTLF is considered [38]. The factors including weather data, characteristics of install devices at areas of interest, history of load and numbers of customers are accounted in it. The factors of economic are taken into account for long period methods of load forecasting.

2.3.

Based on the required time horizons we can further classify the problem into three broad solution domains – heuristic, statistical or econometric, and probabilistic models. Time series datasets have the unique property of dependence on what happened on the past. Therefore, we cannot randomly shuffle the order of the data without affecting the trends. With such dependency simple regression technique for forecasting can randomly show statistical significance even if there is no true correlation, and thus suitable for real-world usage. In the following sections, we will discuss the popular forecasting models used in the electricity demand forecasting domain.

Methods	Time duration	Factors			Uses
		Temperature	Economics	Use of land	
VSTLF	Few minutes	Not compulsory	Not compulsory	Not compulsory	To generate forecasting. Distribution schedule.
STLF	Few hours	Compulsory	Not compulsory	Not compulsory	Maintenance schedule. Allocation of spinning reserve.
MTLF	Days to months	Simulated	Compulsory	Not compulsory	Seasonal forecasting
LTLF	More than year	Simulated	Simulated	compulsory	To plan about growth of generation

Figure 2.3: Different load forecasting methods [source: ..]

2.4   Forecasting Models

Various techniques have been developed for electricity demand forecasting during the past few years. Statistical models are widely adopted for the load forecasting problem, which include linear regression models, stochastic process models, exponential smoothing and ARIMA models. To incorporate the nonlinearity of electricity demand series, artificial neural networks (ANNs) have also received substantial attentions in load forecasting with good performance reported. Neural networks have been shown to have the ability not only to learn the load series but also to model an unspecified nonlinear relationship between load and weather variables. Recently, machine learning techniques and fuzzy logic approaches have also been used for load forecasting or classification and achieved relatively good performances.

Techniques such as Artificial Neural Networks (ANNs) and fuzzy logic models are often used due to their ability to process non-linear patterns in real-time. For very short-term forecasting, methods that can capture rapid fluctuations in demand are essential. Additionally, autoregressive integrated moving average (ARIMA) models and local fuzzy reconstruction can handle time-series data with high frequency, making them suitable for forecasting electricity demand in timeframes of seconds to minutes. These models perform well in scenarios where the primary objective is operational control, such as maintaining grid stability or reacting to sudden changes in demand.

For short-term forecasting, covering hours to a few days, models that integrate external factors such as weather conditions, time of day, and social behaviour become critical. Generalized Additive Models (GAMs) and Support Vector Machines (SVMs) are often applied for their ability to account for non-linear relationships between demand and influencing variables like temperature, day of the week, and holidays. Hybrid models, which combine machine learning techniques with statistical methods like exponential smoothing or seasonal ARIMA, are also effective for this timeframe, as they can capture both the short-term trends and seasonal patterns inherent in electricity consumption. [Wang et al.(2021)Wang, Sun, Chen, Zeng, Kong, Chen, Gu highlights that although these traditional techniques such as fuzzy linear regression, exponential smoothing, Automatic Regressive Moving Average (ARMA) have the advantages of algorithmic simplicity but they are not scalable to handle large dataset

and model complicated relationships. **RNNs** are commonly used for load forecasting when there is a need to model temporal dependencies over shorter sequences. They are effective in predicting short-term fluctuations in electricity demand, especially when there is a clear trend or pattern over recent time intervals, such as hourly or daily demand variations. However, traditional RNNs suffer from the vanishing gradient problem, which limits their ability to retain information over long time sequences. This makes them less effective for capturing long-term dependencies or seasonal patterns in electricity consumption. To overcome the limitations of RNNs, **LSTM networks** are often preferred for load forecasting, especially when dealing with longer time horizons or when it's important to capture both short-term and long-term trends in electricity demand. LSTM networks are designed with a memory cell that can store information for extended periods, allowing the model to retain past information over long time sequences. This makes LSTMs particularly effective for **short-term to medium-term forecasting**, where demand patterns depend on both recent trends (such as hourly or daily fluctuations) and more distant patterns (such as weekly or seasonal cycles). LSTMs are also highly effective when electricity demand is influenced by a range of factors, such as weather conditions, time of day, and past consumption trends.

Medium-term forecasting, which spans weeks to months, focuses on predicting consumption patterns based on broader trends, such as seasonal variations, economic cycles, and energy market dynamics. Time-series models like SARIMA (Seasonal ARIMA) and Exponential Smoothing State Space (ETS) models are widely used for their ability to capture seasonal effects and trends. Additionally, machine learning techniques like Gradient Boosting Machines (GBM) or Random Forests can be leveraged to include more complex interactions between variables, such as the effect of emerging technologies like rooftop solar panels. These models are crucial for tasks such as planning maintenance or scheduling energy generation resources.

For long-term electricity demand forecasting, **econometric models** and **scenario-based models** are often preferred over short-term forecasting methods like neural networks. This is because long-term forecasting, which spans months to years, requires capturing structural trends, external factors, and macroeconomic indicators that influence electricity demand over extended periods.

The simplest form of forecasting is Moving Average (MA) which predicts the future values as the average of n previous values. With a larger window we can smooth out the noise to find the actual tend but the forecasts will start to lag as we were to go back further and further to calculate the average. Exponential Smoothing (ETS) eliminates this problem by exponentially decrease the weights of the averages over time. This prioritises more recent values over the older values, however, as the averaging window size increases the overall trend starts to lag again. Additionally, most real-world time series are not stationary i.e., the change in mean or variance over time are not the same at different observation points. Autoregressive Integrated Moving Average (ARIMA) [Box and Jenkins(1970)] extends this idea further by modelling the dependent relationship between an observation and a set of previously lagged observations by taking the difference of successive values many times to make the time series stationary. This is like saying that it will be hot tomorrow since it has been hot so far this week as the high air temperature had not varied much. In short, the $ARIMA(p, q, d)$ model applies a dth

order differencing on a time series and takes the last p terms in autoregressive way and the include the last q moving average terms to generate the forecast over a specified time horizon. However, the ARIMA model does not consider the repeating patterns within the time series like temperature increases in the daytime while decreases in night or hotter summer and colder winter. Seasonal ARIMA (SARIMA) [Box and Jenkins(1970)] deals with this challenge by taking the values of p, q, and d over a season period m instead of the immediate past. Auto ARIMA [Hyndman and Khandakar(2008)] scales this process by finding the best combination of $< p, q, d >$ automatic way. Typically, ARIMA models and its other variants forecast very well but it's hard to scale for large datasets and often too complex to fine tune without significant domain expertise. The Theta Forecast [Assimakopoulos and Nikolopoulos(2000)] is another well performed model that uses a parameter θ to smooths the local curvature of a time series. In the original implementation two theta lines are used and the final forecast is the average of them. Later, [Hyndman et al.(2002)Hyndman, Allan and Koehler] showed that an ETS with a drift term can perform equally to the original Theta Forecast and since then adopted to most of its later implementations. [Wang et al.(2021)Wang, Sun, Chen, Zeng, Kong highlights that although these traditional techniques such as fuzzy linear regression, exponential smoothing, Automatic Regressive Moving Average (ARMA) have the advantages of algorithmic simplicity but they are not scalable to handle large dataset and model complicated relationships.

Another approach is to first apply Fourier Transformation (FT) to decompose a stationary time series (i.e., removing trends and seasonality) from time domain to frequency domain as temporal frequencies. Once decomposed we can then apply Fast Fourier Transform (FFT) [Shannon(1948)] to filter out the noise and only select the few promising frequencies to perform an Inverse Fast Fourier Transform (IFFT) to reconstruct the time series and add back the trend and seasonality. Due to its unique properties FFT is a versatile and fast generic tool to identify trends and noise filtering in time series and therefore often used as to establish a baseline forecast. A more recent and very popular Generalised Additive Model (GAM) for regression technique is *Prophet* [Taylor and Letham(2017)]. It is designed to optimally handle business forecasting tasks featuring time series captures at the hourly, daily, or weekly level with ideally at least a full year of historical data. It can also handle missing data and outliers, strong seasonality effects occurring daily, weekly, yearly, holidays and other special one-time events that do not necessarily follow the seasonality patterns. These unique properties make it very suitable for both short-term and mid-term electricity forecasting as we will discuss the *Prophet* model in greater details in following subsection. Instead of finding a single point forecast from a time series using statistical methods discussed in the previous subsection, we can also apply probabilistic methods like probabilistic regression or other machine learning models. In the simplest form to cast a time series forecasting as regression we need to first convert the observations i.e., samples as independent and identically distributed. We can treat the set of time lags over a fixed length sliding window as independent feature variables and their succeeding time lag as the corresponding target variable to convert the time series into a dataset for interpolation using regression. We can also use rolling and seasonal rolling time window or Exponentially Weighted Moving Averages (EWMA) instead of sliding window [Roberts(1959)].

This kind of transformation using time lags is called time-delay embedding. Another approach is to embed time into features so that we can train ML models. For instance, timestamps associated with time series can be treated as calendar features i.e., time periodicity like year, month, quarter, day, hour, minutes, seconds, etc. We can also use the timestamps as monotonically increasing integers and treat them collectively as time elapsed feature. tsfeatures (??) and tsfresh (??) are two popular open-source libraries to extract time series features for ML models. One straightforward way is to use Linear regression and estimate using least squares to minimise the Mean Squared Errors (MSE) [Lai et al.(2005)Lai, Yu and Wei]. But multi-collinearity in time series can lead to unstable and unreliable regression models due to the high correlation between variables caused by trends, seasonality, or lagged values. Thus, it's difficult to accurately estimate the coefficients and interpret the relationships between variables. Regularization techniques like applying L1 and L2 norms e.g., lasso and ridge regressions can help address multicollinearity in linear regression by penalising large coefficients. However, we cannot capture the non-linear relationships within the independent features using a linear model. For example, sudden increase in air-conditioning usage after temperature reaches a certain threshold.

A common way to capture such non-linear relationship is to use a Decision Tree (DT) [Spiliotis(2022)] and split branches based on feature importance such as important time lag moment or seasonal rolling. But a DT can easily overfit as it often relies on specific feature importances. We can overcome this by using ensemble learning. Bagging and Radom Forest are two popular ensemble techniques that use a set of DTs as weak learners. Furthermore, we can use Gradient Boosted DT (GBDT) technique for sequential learning and its ability to handle complex relationships for many time series applications. Popular boosting implementations like XGBoost [Chen et al.(2020)Chen, Huang and Zheng] and Light Gradient Boosting Machine (LightGBM) [Salinas et al.(2017)Salinas, Flunkert, Gasthaus, Januschowski and Van den Oord] employ techniques like row and column sampling, and optimised tree growth algorithms to reduce training time for large time series dataset and improve prediction accuracy.

Now, we shift our focus into forecasting from time series using supervised learning models like neural network and deep learning. Transforming time series data for supervised learning models requires considerations of temporal dependencies and trends. Encoder-Decoder and Feed Forward Networks (FNNs) are suitable for fixed-length sequences while Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks can better handle variable-length sequences and long-term dependencies. RNNs address the sequential nature of time series data by incorporating a recurrent connection allowing information from previous observations to be passed to subsequent ones [Hochreiter and Schmidhuber(1997a)]. While RNNs are capable of modelling long-term dependencies they suffer from the vanishing gradient problem. LSTMs, a variant of RNNs, introduces gates that selectively control the flow of information, mitigating the vanishing gradient problem. LSTMs have been successfully applied to a wide range of time series tasks including forecasting, anomaly detection, and classification [Hochreiter and Schmidhuber(1997b)]. Convolutional Neural Networks (CNNs) originally designed for image processing but later have been adapted for time series analysis too. CNNs can effectively capture

9

local patterns and dependencies making them suitable for time series classification and anomaly detection [?]. Lastly, Attention mechanisms and Transformers can capture global dependencies and the importance of different time steps for complex time series tasks. Attention mechanisms can model long-range dependencies in sequence-to-sequence tasks. Attention allows the model to focus on relevant parts of the input sequence when making predictions [?]. Transformers leverage self-attention mechanisms to capture dependencies between different time steps in the input sequence. Transformers have been shown to outperform RNNs and LSTMs in many sequence modelling tasks [?].

2.5 Challenges to Find the Ideal Electricity Load Forecasting Method

In this study we focus on short-term demand forecasting which is an essential instrument in power system planning, operations, and control. Many operating decisions are based on load forecasts such as dispatch scheduling of generating capacity, reliability analysis, and maintenance planning for the generators. Overestimation of electricity demand will cause a conservative operation which may lead to overproduction or excessive energy purchase. For instance, in [AEMO(2022)] AEMO uses a monthly regression model based on five years (60 months) of historical data. The choice of five years data strikes a balance between ensuring that the model considers only relatively recent univariate consumption trends and behaviours while being long enough to capture seasonality and contain enough multivariate observations to be statistically meaningful. Univariate models, while simpler, often struggle to capture complex relationships and seasonal patterns in load data. Multivariate models, incorporating additional factors like temperature, humidity, and economic indicators, can improve accuracy but introduce challenges in data collection, preprocessing, and model complexity, thus require careful feature engineering. Univariate models use only historical demand data to make future predictions, while multivariate models consider other variables such as atmospheric variations and calendars along with historical demand data in the study of STLF [Chapagain et al.(2023)Chapagain, Gurung, Kulthanavit and Kittipiyakul]. [Wang et al.(2016 Chen et al.(2015)Chen, Wang, Liu, Wang and Liu] both highlight the limitations of traditional methods and the growing interest in advanced hybrid techniques. Therefore, the selection of an ideal forecasting method depends on factors such as data availability, computational resources, and the desired level of accuracy, making it a complex and ongoing research area.

[Taylor and McSharry(2009)] develops both univariate and multivariate models and stated that the univariate models had good prediction capability. In univariate time series models, the historical electricity demand data are arranged with correlated past lags to capture the demand patterns even when the data are limited. [McCulloch et al.(2001)McCulloch, Tsay and Wu] obtains improved accuracy by including the temperature as a variable, recognising that weather conditions play a crucial role in forecasting performance. [Fan and Hyndman(2012)] proposes a semi-parametric additive models are proposed to estimate the relationships between demand and the driver variables. Specifically, the inputs for these models

are calendar variables, lagged actual demand observations and historical and forecast temperature traces for one or more sites in the target power system. The proposed methodology has been used to forecast the half-hourly electricity demand for up to seven days ahead for power systems in the Australian National Electricity Market (NEM). To achieve more efficient and accurate load forecasting [Suo et al.(2019)Suo, Song, Dou and Cui] establishes a multi-dimensional short-term load forecasting model based on XGBoost.The decision forest composed of many decision trees is the final learning model of XGBoost. It tries to correct the residual of all the previous weak learners by adding new weak learners. When these learners are combined for final prediction the accuracy will be higher than that of a single learner. The selection of appropriate hyper parameters of XGBoost is directly related to the performance of the model but there is no universal and scientific method to determine the hyper parameters. In order to reduce the randomness and blindness, based on the second search of multi-dimensional grids, a method of hyper parameters optimisation is proposed. Each hyper parameters combination is attempted in a grid traversal manner in turn.

2.6   Performance Metrics/ Model Evaluations

To check the correctness of the methods used for the prediction of real values of load, different criteria are utilized to evaluate the techniques of load forecasting. The research of many researchers based on statistical metrics to optimize the precision of their model, newly developed statistical metrics such as probabilistic load forecasting metrics. Due to wide adaptation and extraordinary academic values in industry, literature on probabilistic forecasting is still in developing phase. The most important static metrics used by researchers are shown in table:

Chapter 3

Material and Methods

3.1 Software

R and Python of course are great software for Data Science. Sometimes, you might want to use `bash` utilities such as `awk` or `sed`.

Of course, to ensure reproducibility, you should use something like `Git` and RMarkdown (or a Jupyter Notebook). Do **not** use Word!

3.2 Description of the Data

How are the data stored? What are the sizes of the data files? How many files? etc.

3.3 Pre-processing Steps

What did you have to do to transform the data so that they become useable?

3.4 Data Cleaning

How did you deal with missing data? etc.

3.5 Assumptions

What assumptions are you making on the data?

3.6 Modelling Methods

Based on the literature review we in this study we explore the two popular model - additive model Prophet by Facebook and another one is the machine learning model XGBoost for short term demand forecasting. According to our EDA, the demand has high variability during peak/off-peak, weekdays, hours of the day and most importantly temperature change.

TBW

CHAPTER 4

Exploratory Data Analysis

This is where you explore your data using histograms, scatterplots, boxplots, numerical summaries, etc.

```python
import numpy as np
np.random.seed(1)
np.random.normal(0.0, 1.0, size=10)
```

```
## array([ 1.62434536, -0.61175641, -0.52817175, -1.07296862,  0.86540763,
##        -2.3015387 ,  1.74481176, -0.7612069 ,  0.3190391 , -0.24937038])
```

CHAPTER 5

Analysis and Results

5.1 XGBoost

Having a very simple model is always good so that you can benchmark any result you would obtain with a more elaborate model.

For example, one can use the linear regression model

$$Y_i = \beta_0 + \beta_1 x_{1i} + \cdots \beta_p x_{pi} + \epsilon_i, \qquad i = 1, \ldots, n.$$

where it is assumed that the ϵ_i's are i.i.d. $N(0, 1)$.

5.2 Facebook Prophet

TBW

CHAPTER 6

Discussion

Put the results you got in the previous chapter in perspective with respect to the problem studied.

CHAPTER 7

Conclusion and Further Issues

What are the main conclusions? What are your recommendations for the "client"? What further analysis could be done in the future?

A figure:

Figure 7.1: A caption

In the text, see Figure 7.1.

References

References

[Adapt NSW(2024)] Adapt NSW, 2024. NSW — AdaptNSW – climatechange.environment.nsw.gov.au. https://www.climatechange.environment.nsw.gov.au/my-region/nsw. [Accessed 21-09-2024].

[AEMO(2022)] AEMO, 2022. Forecasting approach – electricity demand forecasting methodology. https://aemo.com.au/-/media/files/electricity/nem/planning_and_forecasting/nem_esoo/2023/forecasting-approach_electricity-demand-forecasting-methodology_final.pdf. [Accessed 21-09-2024].

[Ahmad et al.(2022)Ahmad, Ghadi, Adnan and Ali] Ahmad, N., Ghadi, Y., Adnan, M., Ali, M., 2022. Load forecasting techniques for power system: Research challenges and survey. IEEE Access 10, 71054–71090. doi:10.1109/ACCESS.2022.3187839.

[Assimakopoulos and Nikolopoulos(2000)] Assimakopoulos, V., Nikolopoulos, K., 2000. The theta model: A decomposition approach to forecasting. International Journal of Forecasting 16, 521–530.

[Box and Jenkins(1970)] Box, G.E.P., Jenkins, G.M., 1970. Time series analysis: Forecasting and control. Holden-Day.

[Chapagain et al.(2023)Chapagain, Gurung, Kulthanavit and Kittipiyakul] Chapagain, K., Gurung, S., Kulthanavit, P., Kittipiyakul, S., 2023. Short-term electricity demand forecasting using deep neural networks: An analysis for thai data. Applied System Innovation 6. URL: https://www.mdpi.com/2571-5577/6/6/100, doi:10.3390/asi6060100.

[Chen et al.(2015)Chen, Wang, Liu, Wang and Liu] Chen, C., Wang, P., Liu, H., Wang, Y., Liu, Z., 2015. Electricity load forecasting: A review. Renewable and Sustainable Energy Reviews 41, 515–527.

[Chen et al.(2020)Chen, Huang and Zheng] Chen, J., Huang, X., Zheng, Y., 2020. A novel hybrid forecasting model based on xgboost and lstm for air quality prediction. Journal of Cleaner Production 277, 123527.

[Fan and Hyndman(2012)] Fan, S., Hyndman, R., 2012. Short-term load forecasting based on a semi-parametric additive model. IEEE Transactions on Power Systems 27, 134–141. doi:https://doi.org/10.1109/tpwrs.2011.2162082.

[Hochreiter and Schmidhuber(1997a)] Hochreiter, S., Schmidhuber, J., 1997a. Long short-term memory. Neural Computation 9, 1735–1780.

[Hochreiter and Schmidhuber(1997b)] Hochreiter, S., Schmidhuber, J., 1997b. Long short-term memory. Neural Computation 9, 1735–1780. doi:10.1162/neco.1997.9.8.1735.

[Hyndman et al.(2002)Hyndman, Allan and Koehler] Hyndman, R.J., Allan, D., Koehler, A.B., 2002. Forecasting time series with exponential smoothing: An r package. Journal of Statistical Software .

[Hyndman and Khandakar(2008)] Hyndman, R.J., Khandakar, Y., 2008. Automatic ARIMA modeling. Journal of Statistical Software 27, 1–22.

[Lai et al.(2005)Lai, Yu and Wei] Lai, K.S., Yu, L., Wei, H.H., 2005. Forecasting stock market prices using neural networks and technical indicators. Neurocomputing 64, 93–105.

[McCulloch et al.(2001)McCulloch, Tsay and Wu] McCulloch, R.E., Tsay, R.S., Wu, P., 2001. Forecasting electricity demand: A comparative study. International Journal of Forecasting 17, 3–18.

[NSW EPA(2021)] NSW EPA, 2021. Energy consumption 2021. URL: https://www.soe.epa.nsw.gov.au/all-themes/human-settlement/energy-consumption-2021. accessed: 2024-09-22.

[R Core Team(2017)] R Core Team, 2017. R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing. Vienna, Austria. URL: https://www.R-project.org/.

[Roberts(1959)] Roberts, D.W., 1959. Prediction intervals for exponential smoothing models. Journal of the Royal Statistical Society. Series B (Methodological) 21, 1–29.

[Salinas et al.(2017)Salinas, Flunkert, Gasthaus, Januschowski and Van den Oord] Salinas, D., Flunkert, V., Gasthaus, J., Januschowski, T., Van den Oord, A., 2017. Deepar: Probabilistic forecasting with autoregressive recurrent neural networks .

[Shannon(1948)] Shannon, C.E., 1948. A mathematical theory of communication. The Bell System Technical Journal 27, 379–423.

[Spiliotis(2022)] Spiliotis, E., 2022. Decision trees for time-series forecasting. Foresight: The International Journal of Applied Forecasting 64, 30–44.

[Suo et al.(2019)Suo, Song, Dou and Cui] Suo, G., Song, L., Dou, Y., Cui, Z., 2019. Multi-dimensional short-term load forecasting based on xgboost and fireworks algorithm, in: 2019 18th International Symposium on Distributed Computing and Applications for Business Engineering and Science (DCABES), IEEE Computer Society, Los Alamitos, CA, USA. pp. 245–248.

[Taylor and McSharry(2009)] Taylor, J.W., McSharry, P.E., 2009. Forecasting electricity demand: A review of methods and approaches. Energy Policy 37, 1279–1292.

[Taylor and Letham(2017)] Taylor, S.J., Letham, B., 2017. Facebook prophet: Forecasting time series with additive models. Journal of the American Statistical Association .

[Wang et al.(2021)Wang, Sun, Chen, Zeng, Kong, Chen, Guo and Wang] Wang, Y., Sun, S., Chen, X., Zeng, X., Kong, Y., Chen, J., Guo, Y., Wang, T., 2021. Short-term load forecasting of industrial customers based on svmd and xgboost. International Journal of Electrical Power & Energy Systems 129, 106830. doi:https://doi.org/10.1016/j.ijepes.2021.106830.

[Wang et al.(2016)Wang, Wang, Xu, Zhang, Liu and Wang] Wang, Y., Wang, L., Xu, L., Zhang, B., Liu, X., Wang, J., 2016. A review of electricity load forecasting techniques. Renewable and Sustainable Energy Reviews 56, 176–189.

[www.soe.epa.nsw.gov.au()] www.soe.epa.nsw.gov.au, . Energy Consumption — NSW State of the Environment — soe.epa.nsw.gov.au. https://www.soe.

epa.nsw.gov.au/all-themes/human-settlement/energy-consumption. [Accessed 21-09-2024].

[www.soe.epa.nsw.gov.au(2021)] www.soe.epa.nsw.gov.au, 2021. Energy Consumption 2021 — NSW State of the Environment — soe.epa.nsw.gov.au. https://www.soe.epa.nsw.gov.au/all-themes/human-settlement/energy-consumption-2021. [Accessed 21-09-2024].

Appendix

Codes

Add you codes here.

Tables

If you have tables, you can add them here.

Use https://www.tablesgenerator.com/markdown_tables to create very simple markdown tables, otherwise use LaTeX.

Tables	Are	Cool
col 1 is	left-aligned	$1200
col 2 is	centered	$12
col 3 is	right-aligned	$1